

**Machine Learning**  
**Assignment 10**  
**Aakanksha Darekar**  
**202200733**  
**09**  
**A1**

---

**Implement a single-layer perceptron using the MNIST dataset.**

**Code:**

```
import numpy as np

import matplotlib.pyplot as plt

from sklearn.metrics import confusion_matrix

import seaborn as sns

from sklearn.preprocessing import LabelBinarizer

# Load MNIST data (using a more reliable method)

def load_mnist():

    from sklearn.datasets import fetch_openml

    mnist = fetch_openml('mnist_784', version=1, as_frame=False)

    X, y = mnist.data, mnist.target.astype(np.uint8)

    return X[:60000], y[:60000], X[60000:], y[60000:] # Standard train-test split

class SingleLayerPerceptron:

    def __init__(self, input_size, output_size, learning_rate=0.01):

        self.weights = np.random.randn(input_size, output_size) * 0.01

        self.bias = np.zeros(output_size)

        self.learning_rate = learning_rate

        self.loss_history = []

        self.accuracy_history = []
```

```
def softmax(self, z): Threshold
```

```
    exp_z = np.exp(z - np.max(z, axis=1, keepdims=True))
```

```
    return exp_z / np.sum(exp_z, axis=1, keepdims=True)
```

```
def forward(self, x):
```

```
    return self.softmax(np.dot(x, self.weights) + self.bias)
```

```
def compute_loss(self, y_true, y_pred): Calculate loss and predict
```

```
    m = y_true.shape[0]
```

```
    return -np.sum(y_true * np.log(y_pred + 1e-10)) / m # Cross-entropy
```

```
def train(self, x, y, epochs=10, batch_size=32):
```

```
    n = x.shape[0]
```

```
    lb = LabelBinarizer()
```

```
    y_onehot = lb.fit_transform(y)
```

```
    plt.figure(figsize=(12, 5))
```

```
    for epoch in range(epochs):
```

```
        # Shuffle and batch
```

```
        indices = np.random.permutation(n)
```

```
        for i in range(0, n, batch_size):
```

```
            batch = indices[i:i+batch_size]
```

```
            x_batch, y_batch = x[batch], y_onehot[batch]
```

```
            # Forward pass
```

```
            preds = self.forward(x_batch)
```

```
            # Backpropagation
```

```
import tensorflow as tf
```

```
from tensorflow.keras.datasets import mnist
```

```
from tensorflow.keras.models import Sequential
```

```
from tensorflow.keras.layers import Dense
```

```
from tensorflow.keras.utils import to_categorical
```

```
# Load dataset
```

```
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
x_train, x_test = x_train.reshape(-1, 784) / 255.0,
```

```
x_test.reshape(-1, 784) / 255.0
```

```
y_train, y_test = to_categorical(y_train), to_categorical(y_test)
```

```
# Build model
```

```
model = Sequential([Dense(10, activation='softmax',  
input_shape=(784,))])
```

```
model.compile(optimizer='sgd', loss='categorical_crossentropy',  
metrics=['accuracy'])
```

```
# Train model
```

```
model.fit(x_train, y_train, epochs=10, batch_size=32,  
validation_data=(x_test, y_test))
```

```
# Evaluate model
```

```
test_loss, test_acc = model.evaluate(x_test, y_test)
```

```
print(f"Test accuracy: {test_acc:.4f}")
```

```
error = preds - y_batch
grad_w = x_batch.T @ error / batch_size
grad_b = np.mean(error, axis=0)
```

if not given

```
# Update weights
self.weights -= self.learning_rate * grad_w
self.bias -= self.learning_rate * grad_b
```

```
# Track progress
train_pred = np.argmax(self.forward(x), axis=1)
acc = np.mean(train_pred == y)
loss = self.compute_loss(y_onehot, self.forward(x))
```

```
self.loss_history.append(loss)
self.accuracy_history.append(acc)
```

```
# Visualize first 5 epochs
if epoch < 5:
    self.visualize_weights(epoch)
```

```
print(f"Epoch {epoch+1}/{epochs} - Loss: {loss:.4f}, Accuracy: {acc:.4f}")
```

```
# Plot training curves
self.plot_training_curves()
```

```
def visualize_weights(self, epoch):
    plt.figure(figsize=(10, 5))
    for i in range(10):
        plt.subplot(2, 5, i+1)
```

```
plt.imshow(self.weights[:, i].reshape(28, 28), cmap='viridis')
plt.title(f"Digit {i}")
plt.axis('off')
plt.suptitle(f"Learned Weight Visualizations (Epoch {epoch+1})")
plt.tight_layout()
plt.show()
```

```
def plot_training_curves(self):
    plt.figure(figsize=(12, 5))
    plt.subplot(1, 2, 1)
    plt.plot(self.loss_history)
    plt.title("Training Loss")
    plt.xlabel("Epoch")
    plt.ylabel("Cross-Entropy Loss")
```

```
    plt.subplot(1, 2, 2)
    plt.plot(self.accuracy_history)
    plt.title("Training Accuracy")
    plt.xlabel("Epoch")
    plt.ylabel("Accuracy")
    plt.tight_layout()
    plt.show()
```

```
def evaluate(self, x_test, y_test):
    y_pred = np.argmax(self.forward(x_test), axis=1)
    acc = np.mean(y_pred == y_test)
    print(f"\nTest Accuracy: {acc:.4f}")
```

```
# Confusion matrix
```

```
cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(10, 8))

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')

plt.title("Confusion Matrix")

plt.xlabel("Predicted")

plt.ylabel("True")

plt.show()
```

```
# Show some example predictions

self.show_examples(x_test, y_test, y_pred)
```

```
def show_examples(self, x, y_true, y_pred, num_examples=10):

    indices = np.random.choice(len(x), num_examples)

    plt.figure(figsize=(15, 3))

    for i, idx in enumerate(indices):

        plt.subplot(1, num_examples, i+1)

        plt.imshow(x[idx].reshape(28, 28), cmap='gray')

        color = 'green' if y_pred[idx] == y_true[idx] else 'red'

        plt.title(f"P: {y_pred[idx]} \nT: {y_true[idx]}", color=color)

        plt.axis('off')

    plt.suptitle("Example Predictions (Green=Correct, Red=Wrong)")

    plt.tight_layout()

    plt.show()
```

```
# Main execution
```

```
if __name__ == "__main__":
```

```
    # Load data
```

```
    X_train, y_train, X_test, y_test = load_mnist()
```

```
# Normalize pixel values

X_train = X_train / 255.0
X_test = X_test / 255.0


# Initialize and train perceptron

perceptron = SingleLayerPerceptron(input_size=784, output_size=10, learning_rate=0.1)

print("Starting training...")

perceptron.train(X_train, y_train, epochs=20, batch_size=128)

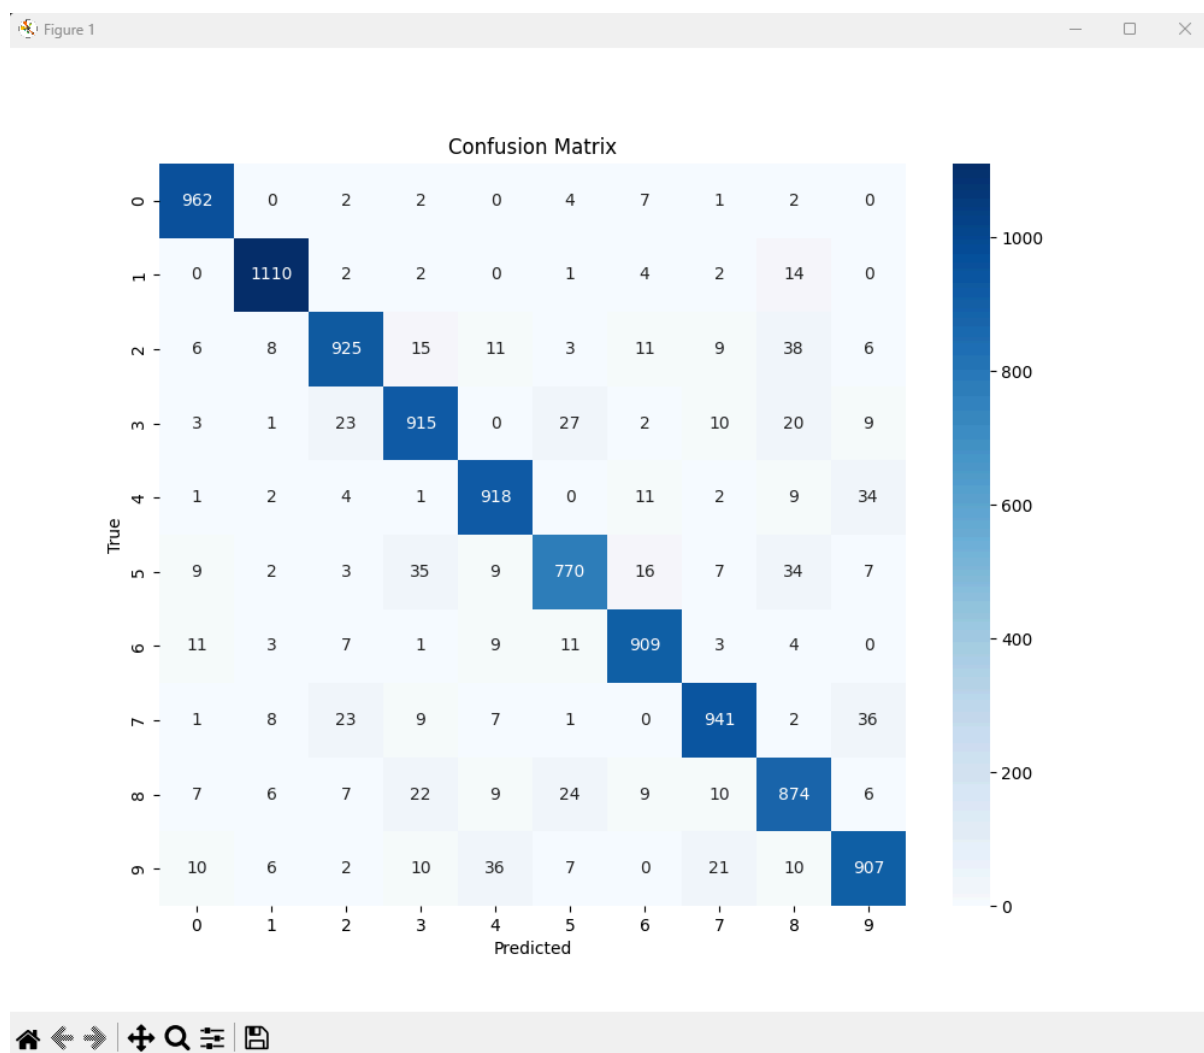
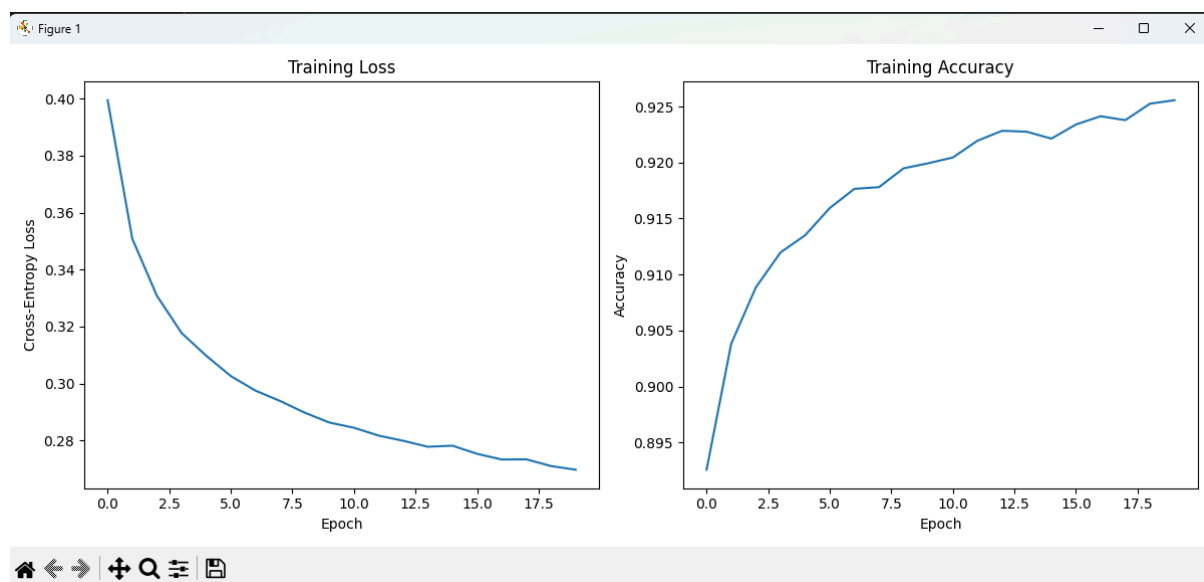

# Evaluate

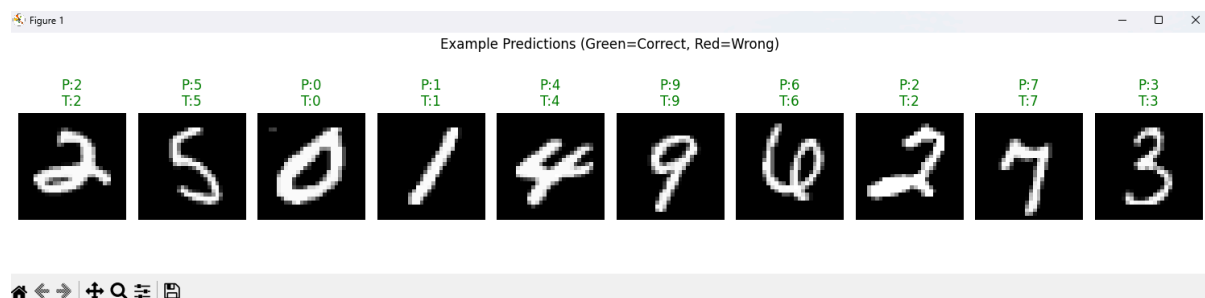
perceptron.evaluate(X_test, y_test)
```

### Output:

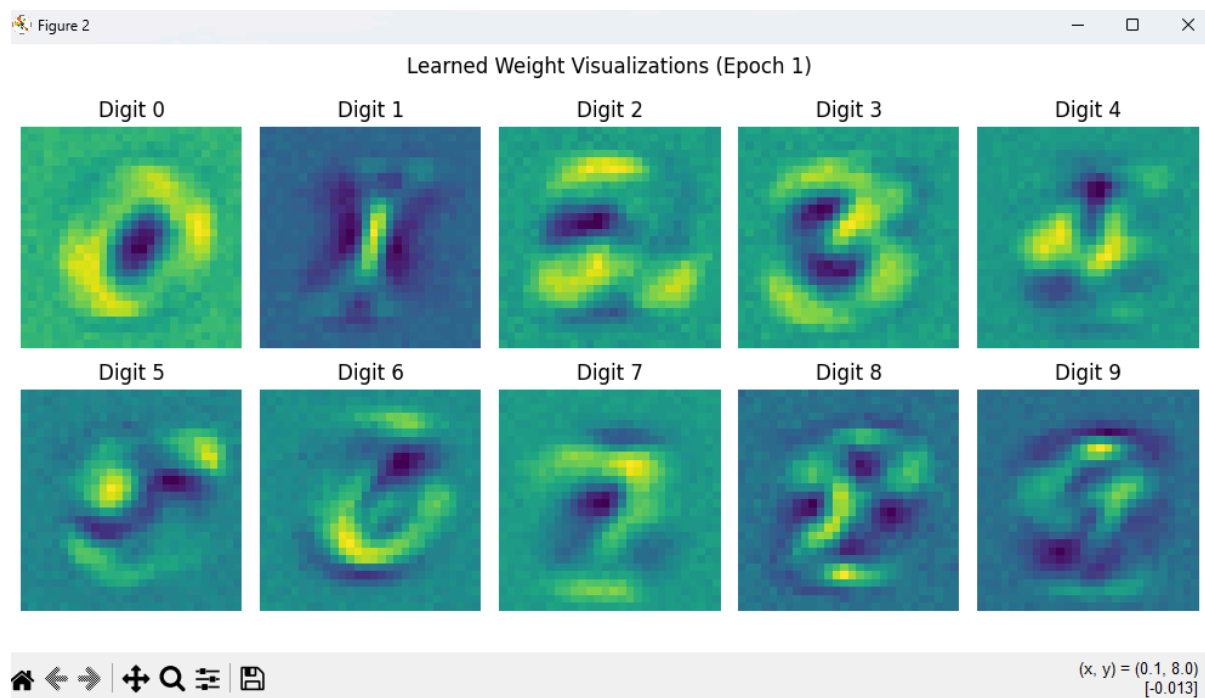
```
Epoch 3/20 - Loss: 0.3304, Accuracy: 0.9085
Epoch 4/20 - Loss: 0.3178, Accuracy: 0.9120
Epoch 5/20 - Loss: 0.3086, Accuracy: 0.9144
Epoch 6/20 - Loss: 0.3025, Accuracy: 0.9159
Epoch 7/20 - Loss: 0.2970, Accuracy: 0.9179
Epoch 8/20 - Loss: 0.2940, Accuracy: 0.9184
Epoch 9/20 - Loss: 0.2906, Accuracy: 0.9193
Epoch 10/20 - Loss: 0.2862, Accuracy: 0.9200
Epoch 11/20 - Loss: 0.2840, Accuracy: 0.9209
Epoch 12/20 - Loss: 0.2816, Accuracy: 0.9217
Epoch 13/20 - Loss: 0.2796, Accuracy: 0.9230
Epoch 14/20 - Loss: 0.2781, Accuracy: 0.9221
Epoch 15/20 - Loss: 0.2763, Accuracy: 0.9233
Epoch 16/20 - Loss: 0.2747, Accuracy: 0.9235
Epoch 17/20 - Loss: 0.2748, Accuracy: 0.9237
Epoch 18/20 - Loss: 0.2723, Accuracy: 0.9250
Epoch 19/20 - Loss: 0.2719, Accuracy: 0.9252
Epoch 20/20 - Loss: 0.2704, Accuracy: 0.9257

Test Accuracy: 0.9241
PS C:\Users\sanke\OneDrive\Documents\Project\Python\Assignments> █
```



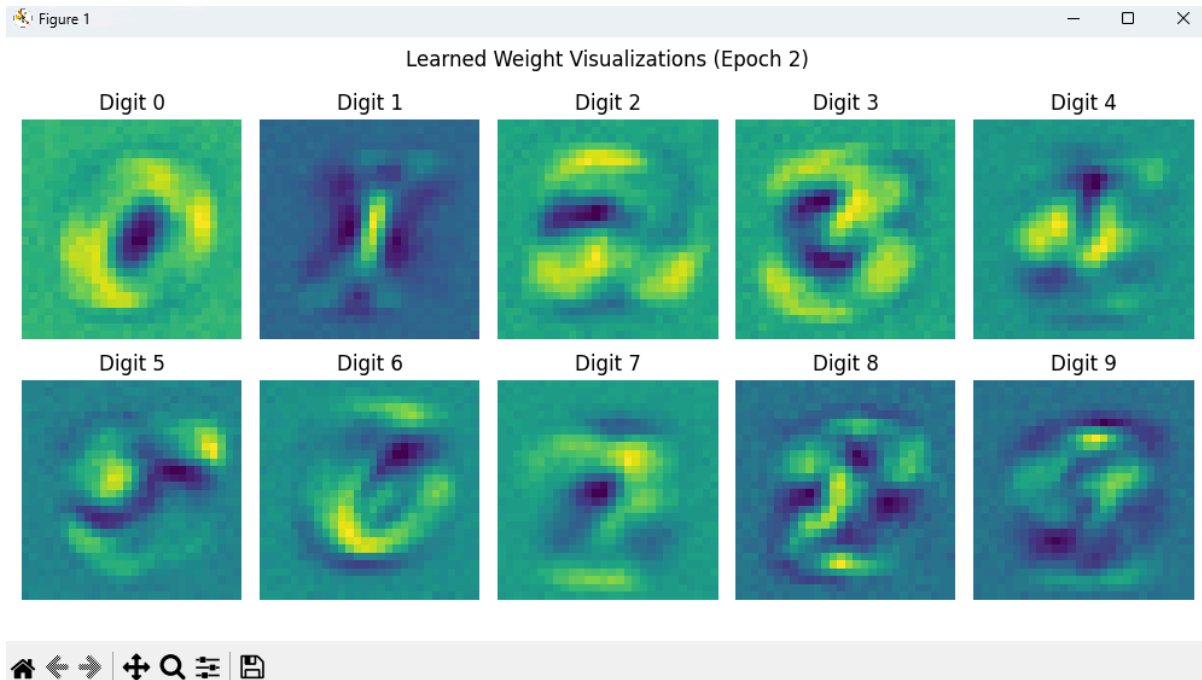


## Epoch 1:



## Epoch 2:





### Epoch 3:

